

# Queues in PostgreSQL

PGDay Australia 2017  
Thomas Munro



# Who am I?

- EnterpriseDB Database Server team member
- PostgreSQL contributor
- First PostgreSQL patch was **SKIP LOCKED** in release 9.5 (topic of this talk)
- Currently working on parallelism and new storage formats

What's a Queue?

Why Put One in an RDBMS?

Example Use Cases

Implementation

Problems

What Could We Do Better?

# queue /kjuː/

noun


1. *Chiefly British* A line or sequence of people or vehicles awaiting their turn to be attended to or to proceed.



# queue /kjuː/

noun

2. *Computing* A list of data items, commands, etc., stored so as to be retrievable in a **definite order**, usually the order of insertion.





# Informal Taxonomy

- Queues
  1. FIFO: First-in-first-out queues
  2. Priority queues
- “Queues”
  3. Specialised queues (merging, reordering)
  4. Unordered/approximately ordered queues


# 1. FIFO Queues

- The order most people think of when they hear the word “queue”
- Often used in low level code because the implementation is simple and fast:  
physical layout reflects logical ordering



# 2. Priority Queues

- Sometimes a different explicit logical order is needed
- Implementation techniques include sets of FIFO queues, **trees** and other data structures associated with sorting





# 3. Specialised “Queues”

- Sometimes we use the word queue more loosely to describe something that gives up strict logical ordering to meet some other goal
- Operating system IO schedulers and elevators/lifts allegedly improve global efficiency by merging and reordering queued requests



# 4. Unordered & Approximately Ordered “Queues”

- Sometimes we don't care about the order that items are retrieved in at all, we just want to process them as quickly as possible
- ... but usually we want at least approximate time ordering for fairness (no arbitrarily stuck messages), but don't need strict global ordering for correctness
- Transactional and concurrent systems blur the order of both insertion and retrieval

What's a Queue?

Why Put One in an RDBMS?

Example Use Cases

Implementation

Problems

What Could We Do Better?

“Meh, why not use  
RabbitMQ/Redis/PGQ/  
*<thing>*?”



# You might consider using a plain old database if...

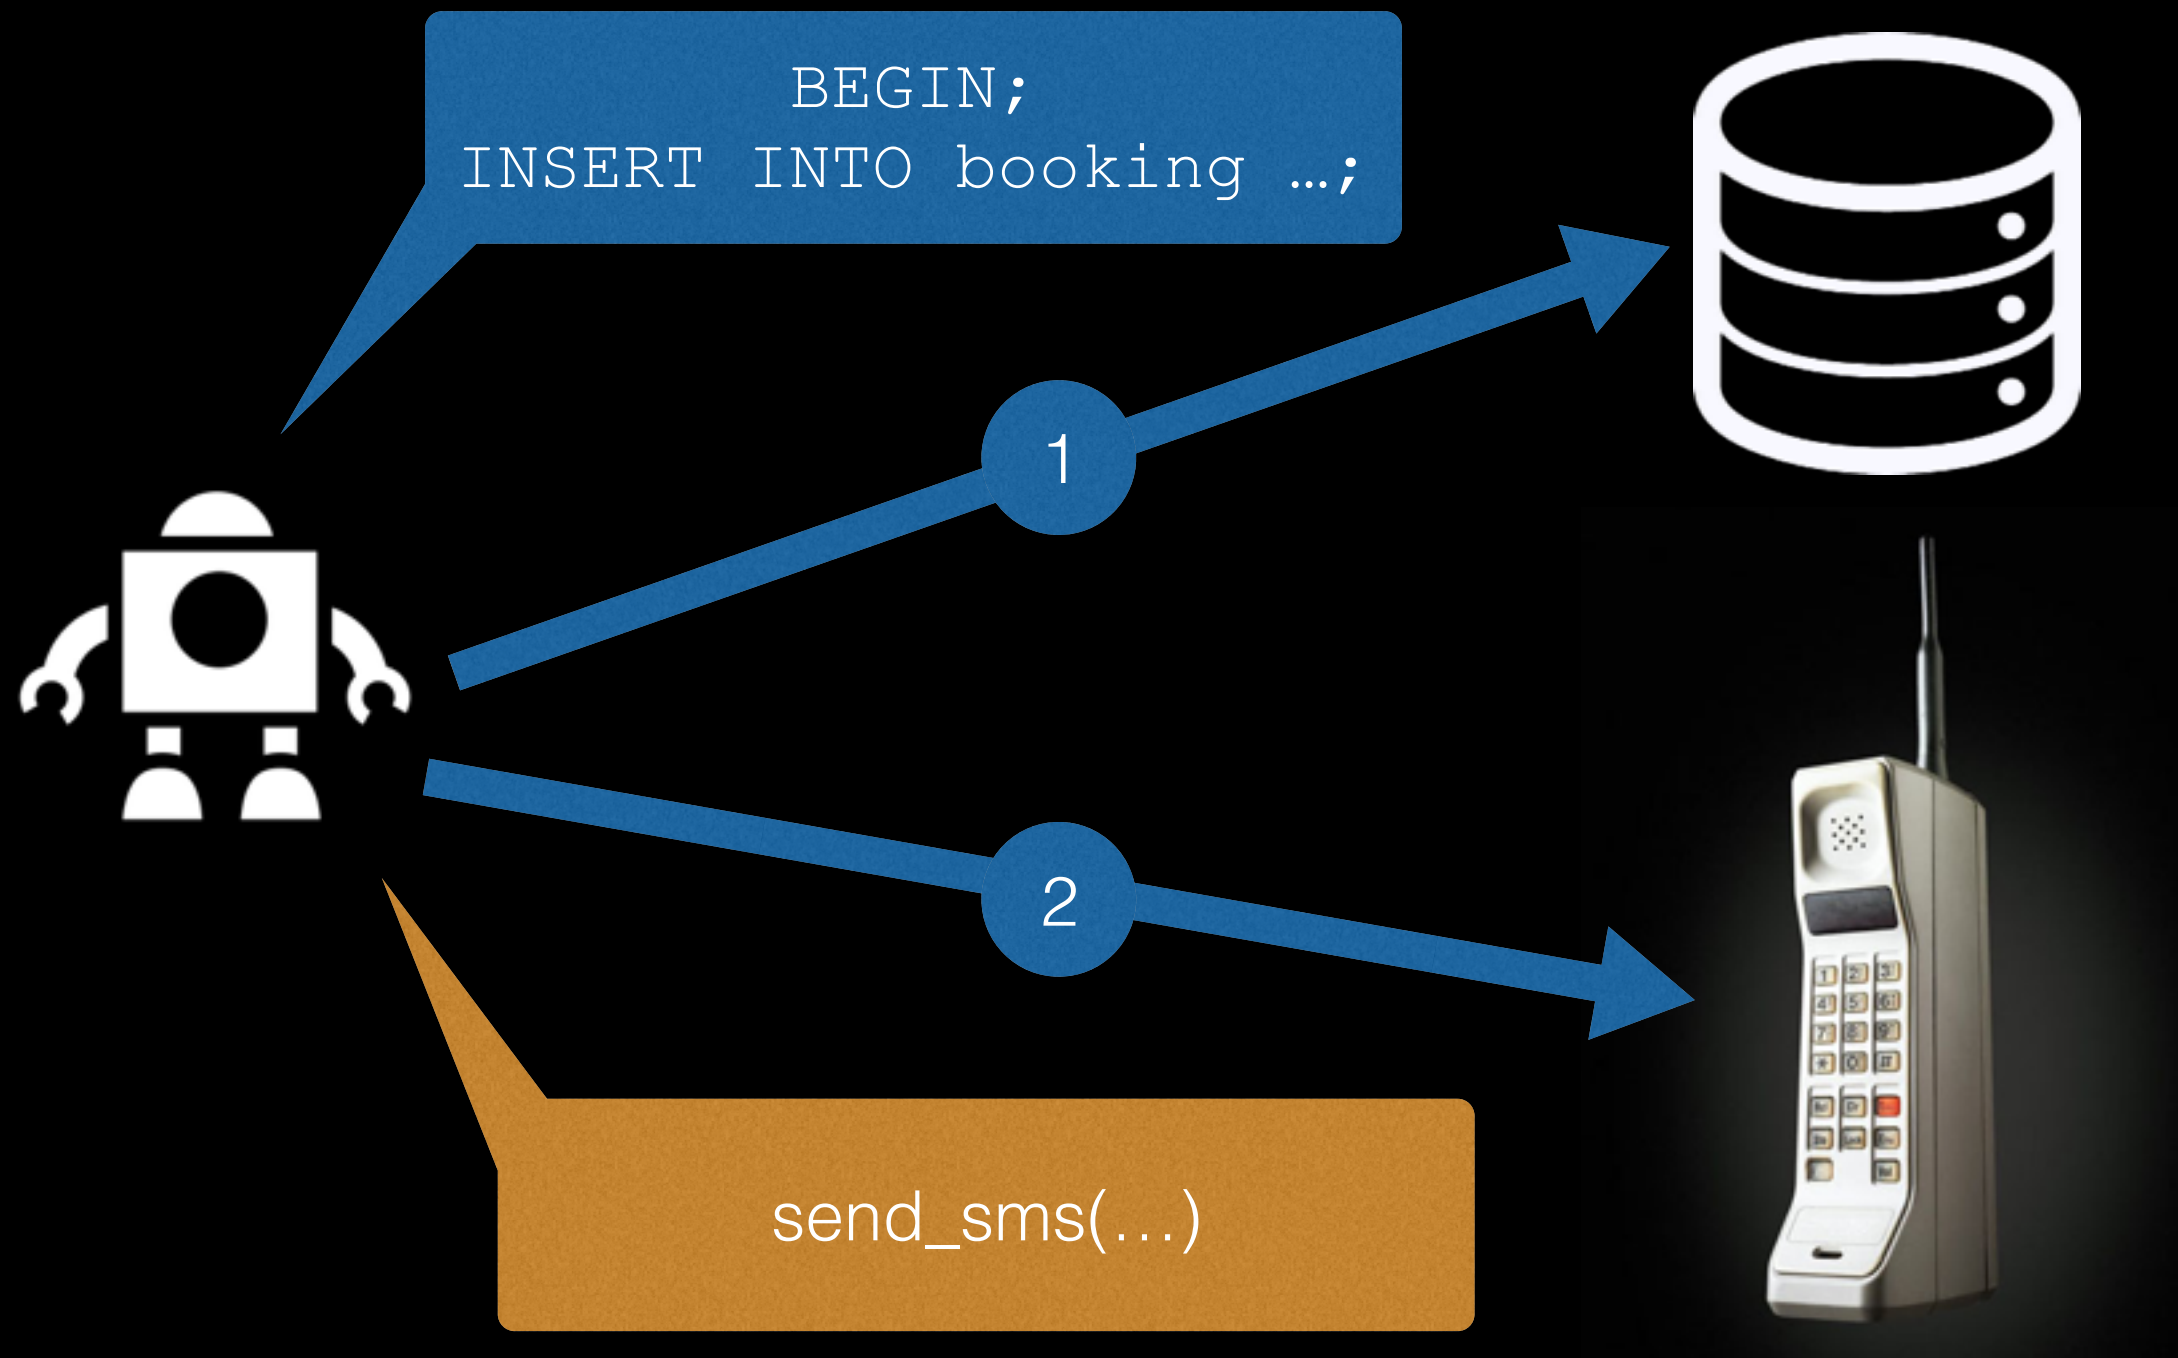
- ... you want reliable persistent message processing that is atomic with respect to other database work (without the complications of distributed transactions)
- ... you don't want the maintenance, backups, failover and risks of new moving parts (message broker daemons)
- ... your message rates and number of consumers are in the range that PostgreSQL and your hardware can handle
- ... you like PostgreSQL enough to attend a conference

What's a Queue?  
Why Put One in an RDBMS?  
Example Use Cases  
Implementation  
Problems  
What Could We Do Better?

# Mixing Transactions with External Effects

- We want to book a seat on a plane
- We also want to send an SMS message with confirmation of the booking and seat number

# Mixing Transactions with External Effects: Take 1

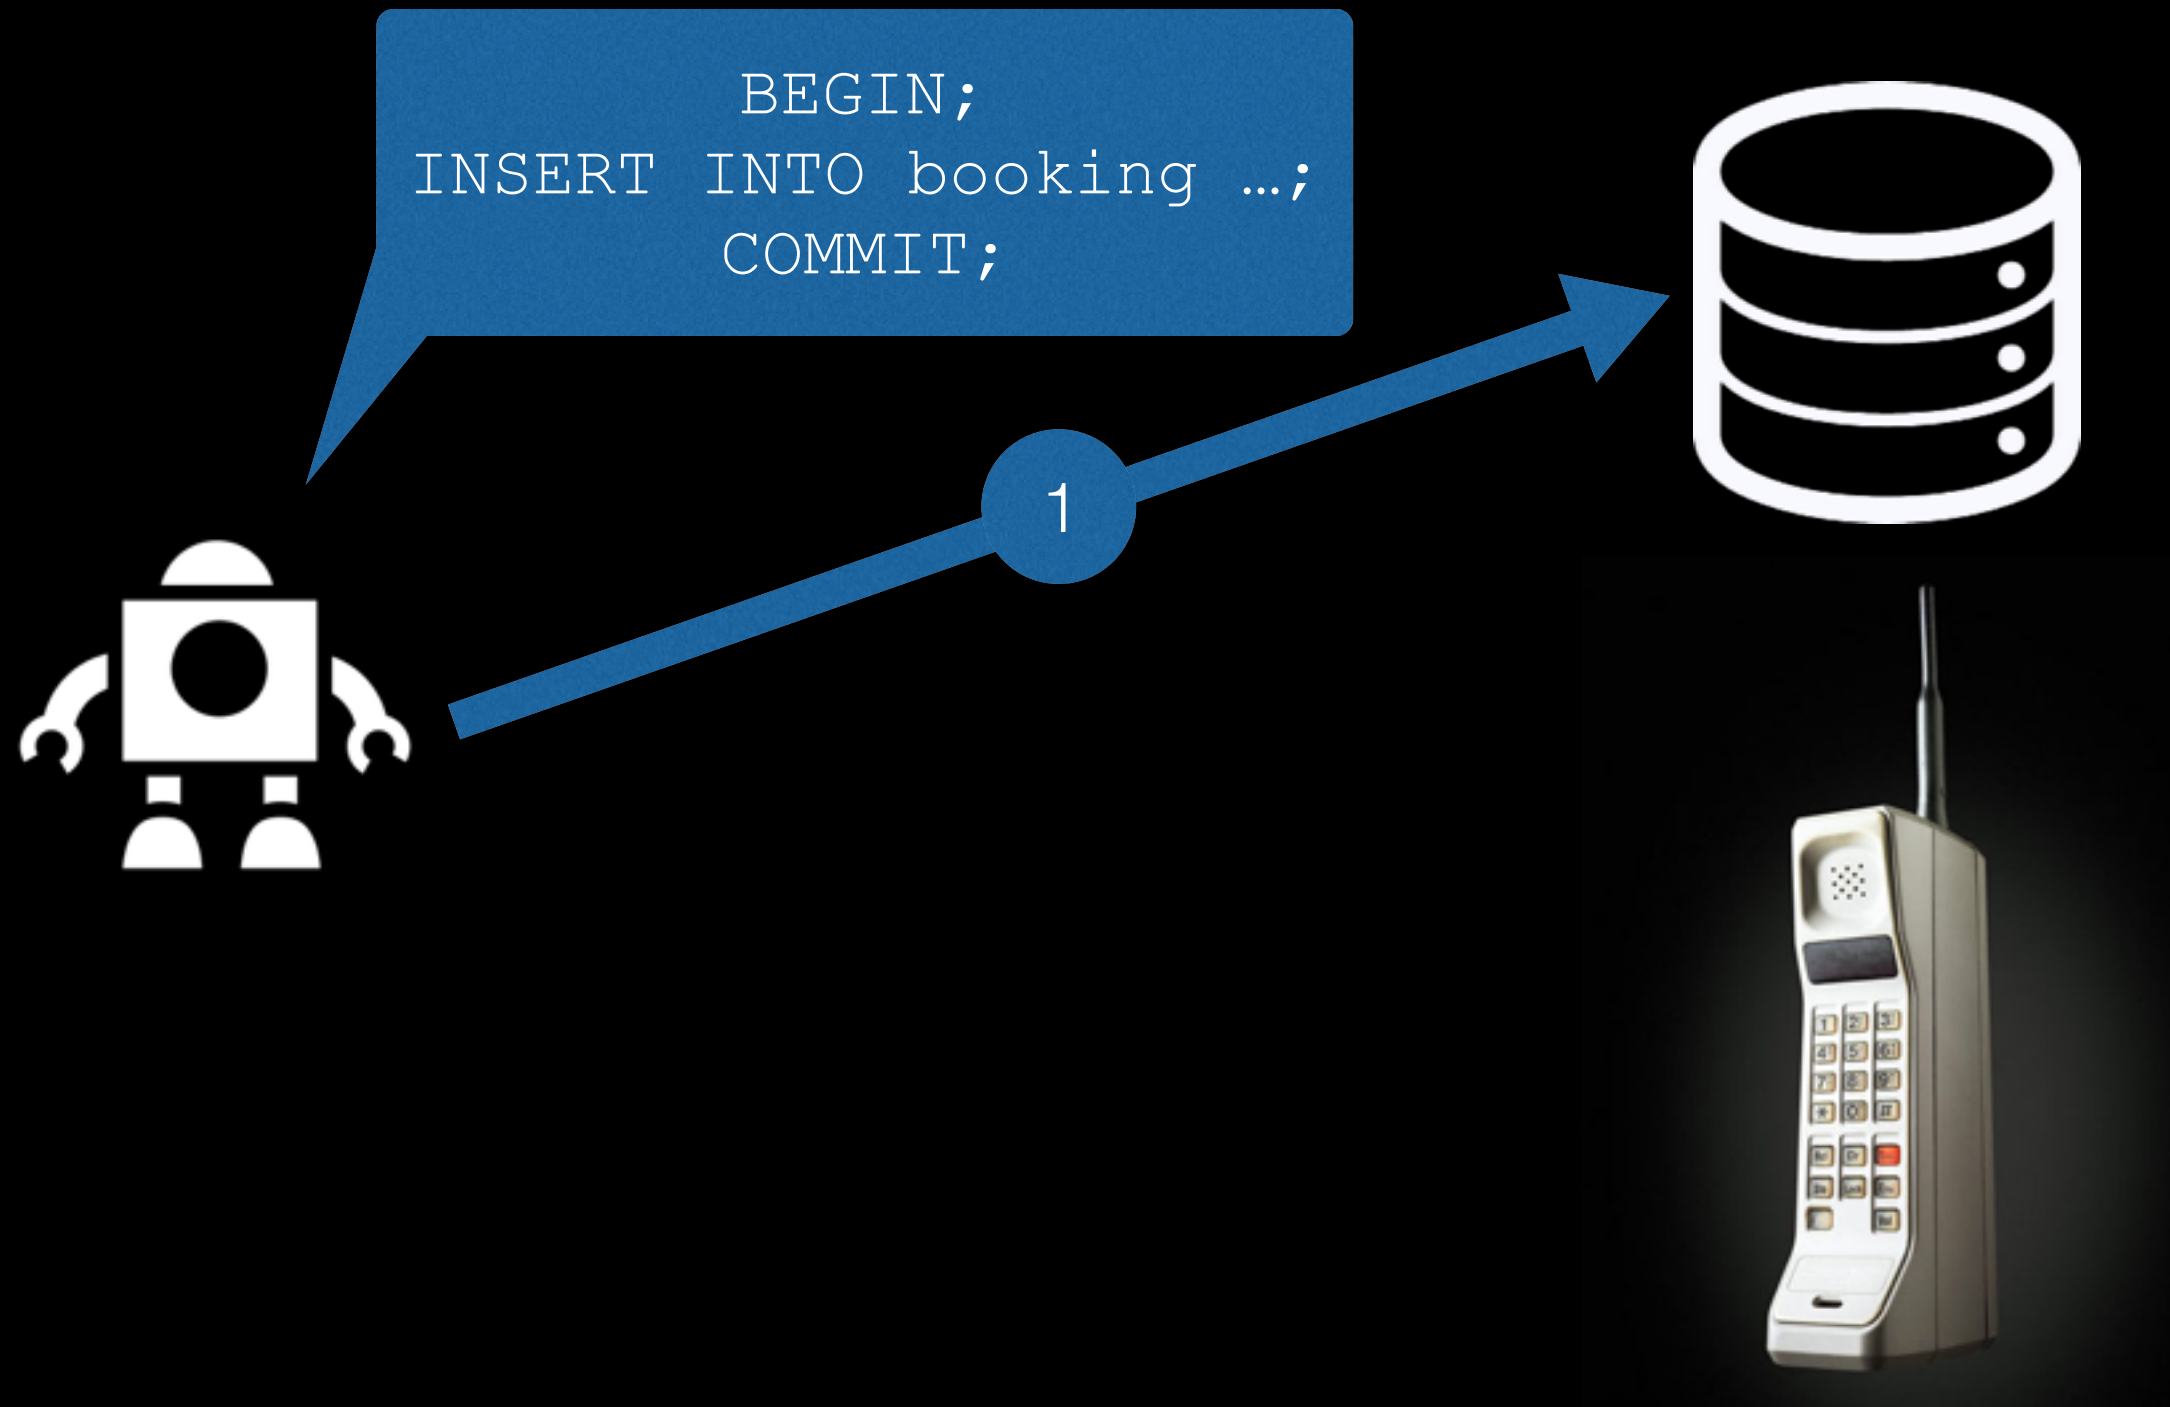




# Mixing Transactions with External Effects: Take 1

Oops: we have sent an SMS but  
forgot the fact it represents due to  
an asteroid/bug/hardware failure  
before COMMIT

# Mixing Transactions with External Effects: Take 2

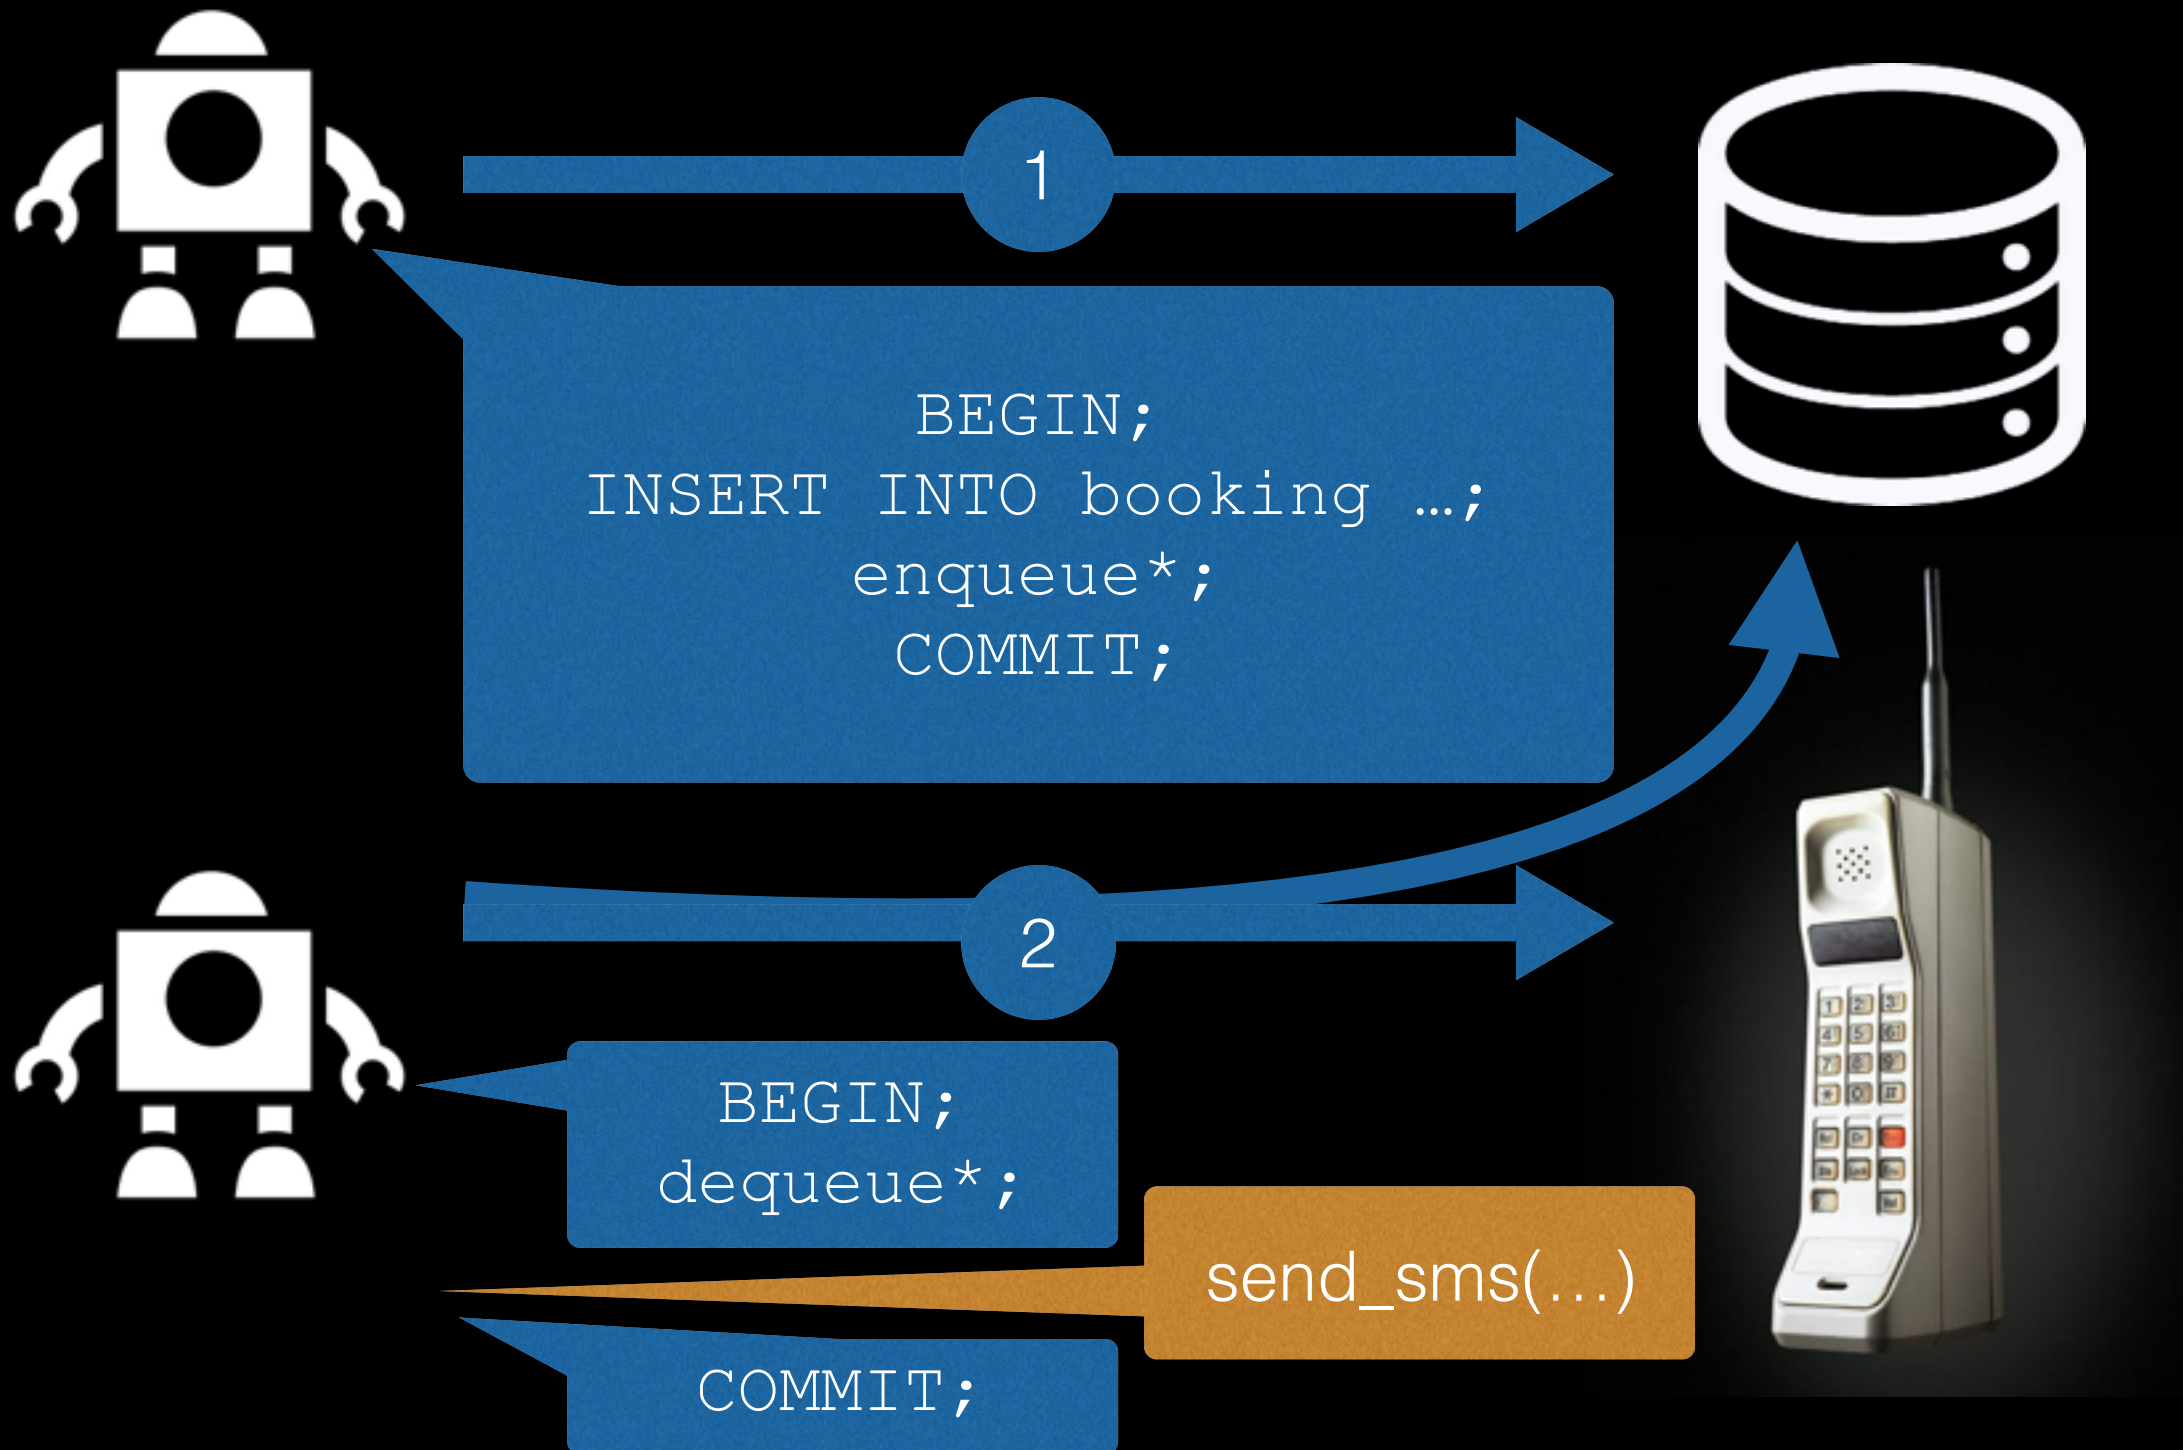




# Mixing Transactions with External Effects: Take 2

Oops: we have committed the fact, but failed to send an SMS due to flood/transient network failure/SMS provider downtime

# Mixing Transactions with External Effects: Take 3







# Mixing Transactions with External Effects

- We establish a new fact (the booking) and record our intention to notify the customer (the entry in the SMS queue) **atomically**
- We remove the queued item after sending successfully (and probably have a retry system if the SMS service is temporarily failing)
- The SMS sending operation should ideally be **idempotent** so that if we fail after sending but before committing the dequeue operation, sending the same message again won't be problematic

# Distributed Computing

- Job control for farming out expensive external computation to worker processes
- Job control for database aggregation work moved out of interactive transactions



What's a Queue?  
Why Put One in an RDBMS?  
Example Use Cases  
Implementation  
Problems  
What Could We Do Better?

# Ingredients

- Messages: Rows in plain old tables
- Priority ordering: ORDER BY
- Signalling: NOTIFY & LISTEN
- Concurrency:
  - None, course grained locking or SERIALIZABLE
  - ... or explicit fine grained locking



# No Physical FIFO

- The relational model (and therefore its approximate earthly embodiment SQL) doesn't expose details of physical ordering or insertion order to the user
- Ordering will therefore need to be a function of values in records supplied at `INSERT` time, and explicitly requested with `ORDER BY` when they are retrieved (it's always a "priority queue"), or unordered

# Enqueue Protocol

- `BEGIN;`  
– any other work  
`INSERT INTO sms_queue (...)`  
`VALUES (...);`  
`NOTIFY sms_queue_broadcast;`  
`COMMIT;`
- Note: if inserting transactions overlap, then it is difficult to generate a key that increases monotonically with respect to commit/transaction visibility order!


# Dequeue Protocol: Take 1

- `LISTEN sms_queue_broadcast;`
- `BEGIN;`  
`SELECT message_uuid, destination, body`  
`FROM sms_queue`  
`ORDER BY insert_time`  
`LIMIT 1;`  
– if found, do something (internal or  
– external + idempotent) and then:  
`DELETE FROM sms_queue`  
`WHERE message_uuid = $1;`  
`COMMIT;`
- – repeat previous step until nothing found
- – wait for notifications before repeating

# Deque Protocol: Take 1

- At isolation levels below `SERIALIZABLE`, this protocol won't work correctly if there are concurrent sessions dequeuing
- At `SERIALIZABLE` level, at most one such overlapping session can succeed (worst case workload for `SERIALIZABLE`)


# Deque Protocol: Take 1




# Dequeue Protocol: Take 2

- `LISTEN sms_queue_broadcast;`
- `BEGIN;`  
`SELECT message_uuid, destination, body`  
`FROM sms_queue`  
`FOR UPDATE`  
`ORDER BY insert_time`  
`LIMIT 1;`  
`– if found, do something (internal or`  
`– external + idempotent) and then:`  
`DELETE FROM sms_queue`  
`WHERE message_uuid = $1;`  
`COMMIT;`
- `– repeat previous step until nothing found`
- `– wait for notifications before repeating`

# Deque Protocol: Take 2



# Deque Protocol: Take 2






# Dequeue Protocol: Take 3


- `LISTEN sms_queue_broadcast;`
- `BEGIN;`  
`SELECT message_uuid, destination, body`  
`FROM sms_queue`  
`FOR UPDATE SKIP LOCKED`  
`ORDER BY insert_time`  
`LIMIT 1;`  
`– if found, do something (internal or`  
`– external + idempotent) and then:`  
`DELETE FROM sms_queue`  
`WHERE message_uuid = $1;`  
`COMMIT;`
- `– repeat previous step until nothing found`
- `– wait for notifications before repeating`

In PostgreSQL 9.4 and earlier which don't have `SKIP LOCKED`, use `pg_try_advisory_lock(x)` in the `WHERE` clause, where `x` is somehow derived from the message ID

# Deque Protocol: Take 3



# Deque Protocol: Take 3



# Dequeue Protocol: Take 3

- The ORDER BY clause is still controlling the time we **start** processing each item, but no longer controlling the order we commit
- Dequeueing transactions that roll back cause further perturbation of the processing order
- Looser ordering is good for concurrency while still approximately fair to all messages
- Stricter ordering is needed for some replication-like workloads with a semantic dependency between messages

What's a Queue?

SQL

Example Use Cases

Implementation

Problems

What Could We Do Better?

# Problems

- Transient failures can be managed with retry counters and two-transaction protocol
- Watch out for ID space running out (32 bit integers)
- If using a SEQUENCE to generate a strict order, be careful of cycling and be aware of behaviour when transactions overlap
- Btrees not correlated with insert/delete order can develop a lot of bloat in high churn tables


# Problems

- Statistics for volatile tables might cause trouble (CF DB2 VOLATILE)
- If there is no ordering requirement at all, in theory you might not even need an index on a queue table (you could use ctid to refer to arbitrarily selected locked rows)
- Default vacuum settings may be insufficient, depending on your workload, leading to bloat and unstable performance







# Vacuuuming



# Vacuuming



# Vacuuming



What's a Queue?

SQL

Example Use Cases


Performance

Problems

What Could We Do Better?

# Undo

- Undo-log based MVCC would provide continuous recycling of space, avoiding bloat and giving smoother performance
- Key question: where should we keep old and deleted versions of rows, so that concurrent transactions that still need them can find them?
- Watch this space — active development area



# Notifications

- It would be nice to have a new wait/notify feature that could handle 'broadcast' like NOTIFY, but also 'notify one': to avoid stampedes of otherwise idle recipients when only one item has been enqueued
- It might be better to do that with a blocking 'wait' function rather than the NOTIFY asynchronous message approach (?)

# Serializable

- Queue-like workloads are the worst case for SERIALIZABLE
- The executor could in theory consider returning tuples in a different order when there is a LIMIT, no [complete] ORDER BY, and another transaction has SIREAD locks on a tuple being returned
- Perhaps this could reduce conflicts in such workloads, allowing higher throughput without giving up the benefits of SERIALIZABLE



<EOF>