

# Need is the mother of all Inventions. It is very true for PostgreSQL 10 Logical Replication

PostgreSQL 10 Logical Replication

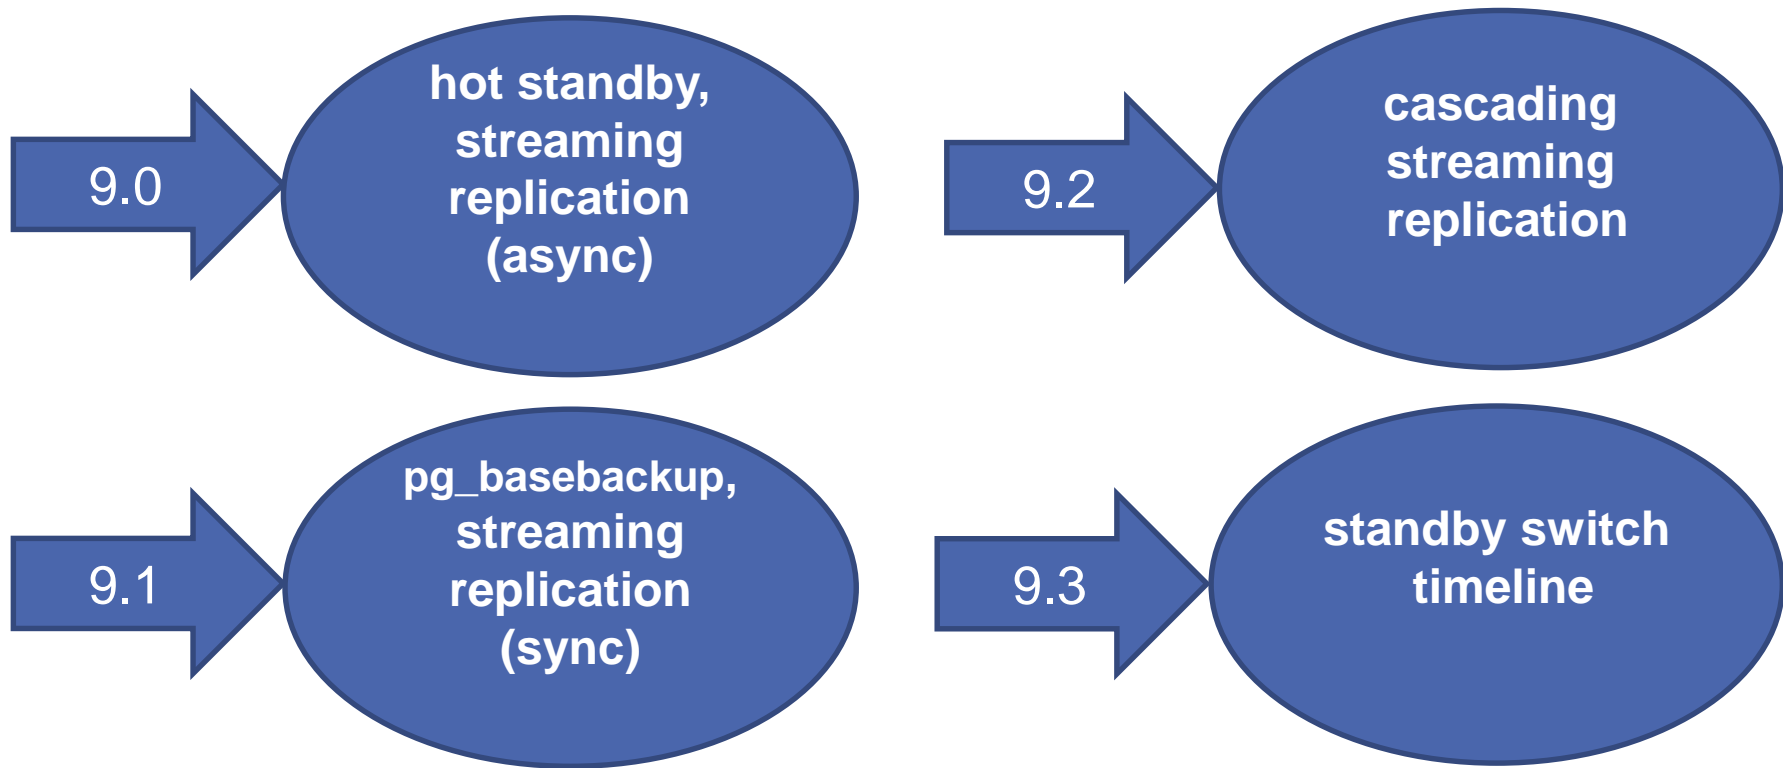
By:  
Rajni Baliyan

# Agenda:





- PostgreSQL replication history.
- Logical Replication: Overview
- Use cases-SR vs LR
- Terminologies used in LR.
  - Publication, Publisher
  - Subscription, subscriber
  - Replication Slot
- Architecture
- Security
- Configuration Settings
- Quick Setup
- Monitoring
- Precautions- Conflicts
- Restrictions/Limitations
- Q&A



# Overview of PostgreSQL replication 9.0 onwards:



# Overview of PostgreSQL replication 9.0 onwards cont..



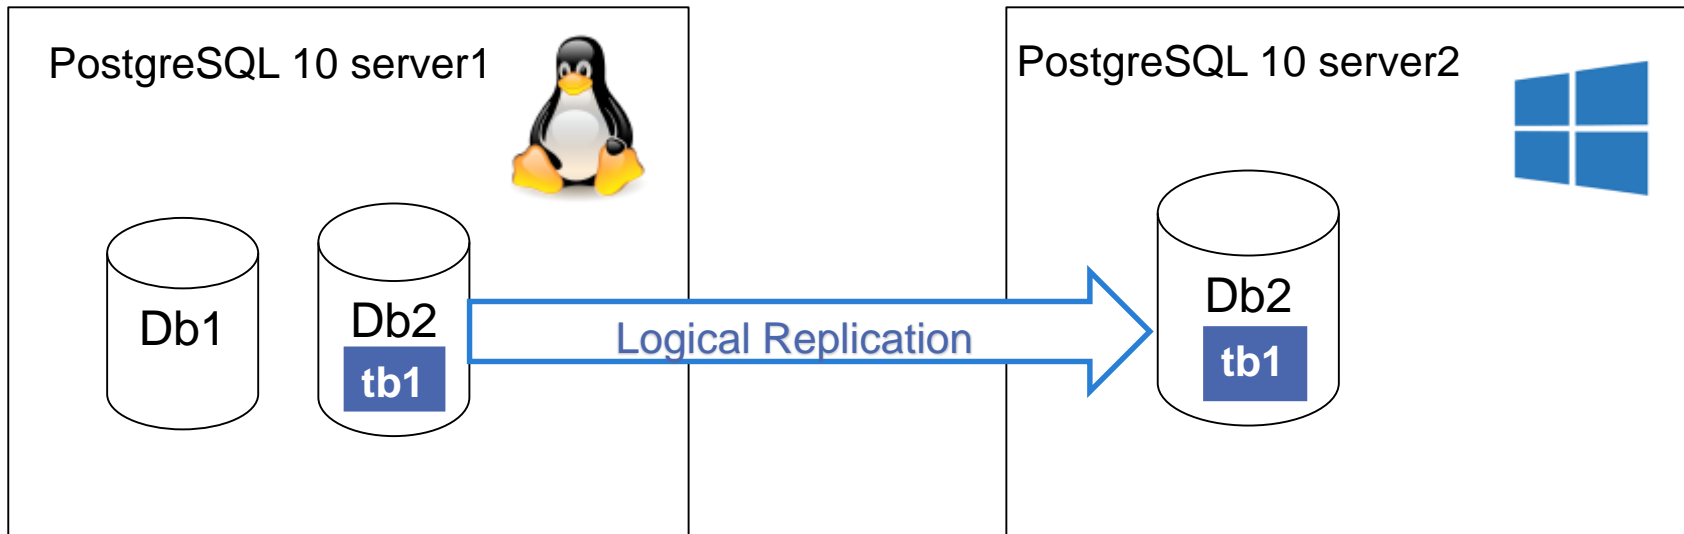
- Replicating/capturing logical changes based upon their replica identity using replication slots.
- Propagate any changes (Insert, Update, Delete, all, combination) to replica.
- Uses a *publish* and *subscribe* model.
  - Publisher-Sender
  - Subscriber- Receiver
- Subscribers pulls data from publications.
- Generally 'asynchronous' in nature.

# Why Logical Replication over Streaming Replication?

Features	Streaming Replication	Logical Replication
<b>Replication Type</b>	Binary replication(byte-by-byte)	Selective row level changes (logical change)
<b>Hardware/OS</b>	Same (Linux-Linux)	Can be different(Linux-Windows)
<b>Replication between major versions</b>	Not supported	Supported.
<b>Replication Level</b>	Instance level	Object level(table)
<b>Consolidation</b>	No consolidation	Consolidating is possible
<b>Replica/ Subscriber- Open Mode</b>	Standby is in read-only mode	Subscriber can be used for write operations

# Use Case1: Cross platform replication

## ➤ Cross Platform : Linux → Windows





# Create Subs. on Windows using Linux pub.

```
pub=# \dRp
```

```
List of publications
```

Name	Owner	All tables	Inserts	Updates	Deletes
testpub	postgres	f	t	t	t

(1 row)

```
logrep=# select slot_name from pg_replication_slots;  
slot_name
```

```
-----  
sync_rep  
testsub  
mysub1  
(3 rows)
```

```
logrep=# select * from emp;  
empno | empname  
-----+-----  
21 | Tom  
22 | Dick  
23 | Harry  
24 | Lee  
(4 rows)
```

Administrator: Command Prompt - psql -U postgres

```
logirep=# select * from emp;  
empno | empname  
-----+-----  
(0 rows)
```

```
logirep=#
```

```
logirep=#
```

```
logirep=# CREATE SUBSCRIPTION mysub1 CONNECTION  
logirep=# 'host=192.168.56.109 dbname=logrep user=postgres'  
logirep=# PUBLICATION testpub;
```

```
NOTICE: created replication slot "mysub1" on publisher  
CREATE SUBSCRIPTION
```

```
logirep=# \dRs
```

```
List of subscriptions  
Name | Owner | Enabled | Publication  
-----+-----+-----+-----  
mysub1 | postgres | t | {testpub}  
(1 row)
```

```
logirep=# select * from emp;  
empno | empname  
-----+-----  
21 | Tom  
22 | Dick  
23 | Harry  
24 | Lee  
(4 rows)
```

## Insert on server1

```
logrep=# select * from emp;
 empno | empname
-----+-----
      21 | Tom
      22 | Dick
      23 | Harry
      24 | Lee
(4 rows)
```



```
logrep=# insert into emp values(25,'Adam');
INSERT 0 1
logrep=# select * from emp;
 empno | empname
-----+-----
      21 | Tom
      22 | Dick
      23 | Harry
      24 | Lee
      25 | Adam
(5 rows)
```

## Select on server2

Administrator: Command Prompt - psql -U postgres

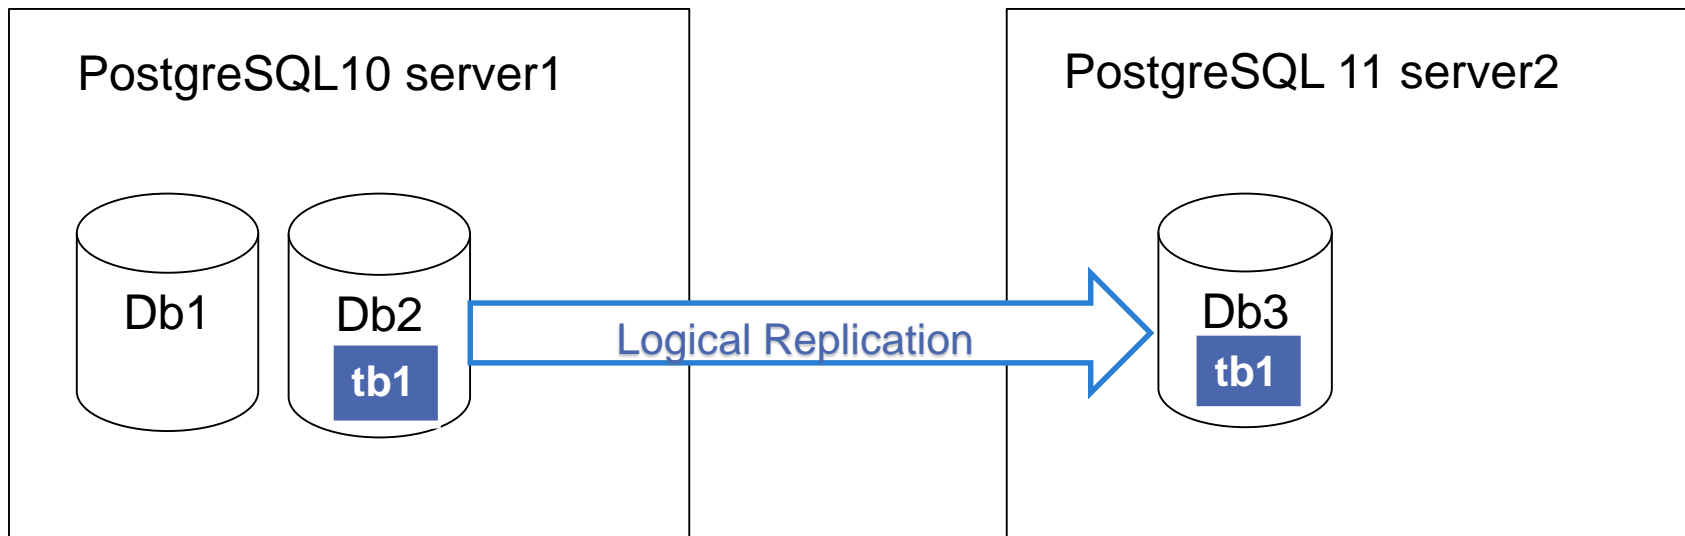
```
logirep=# select * from emp;
 empno | empname
-----+-----
      21 | Tom
      22 | Dick
      23 | Harry
      24 | Lee
(4 rows)
```



```
logirep=# select * from emp;
 empno | empname
-----+-----
      21 | Tom
      22 | Dick
      23 | Harry
      24 | Lee
      25 | Adam
(5 rows)
```

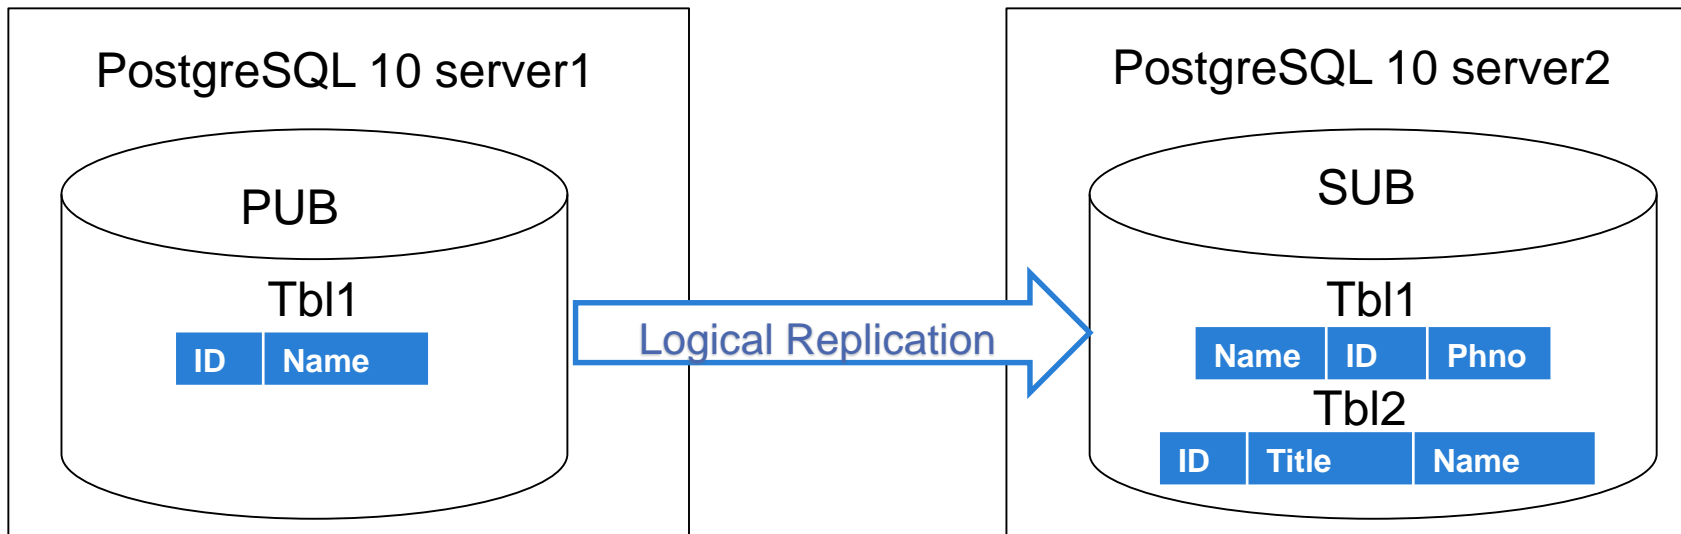
## Use Case2: Cross version:

- **Cross Version: Replication between different major versions of PostgreSQL**



# Use Case3: Write operation @Subscriber:


- **Modifying objects and adding new tables, insert records on subscriber.**




# Use Case4: Consolidating multiple databases

- Can be helpful for reporting or analytical purposes
- Can replicate even few tables from one DB to another on different servers.


## ■ PostgreSQL Server1



## ■ PostgreSQL Server2(reporting)



## ■ PostgreSQL Server3

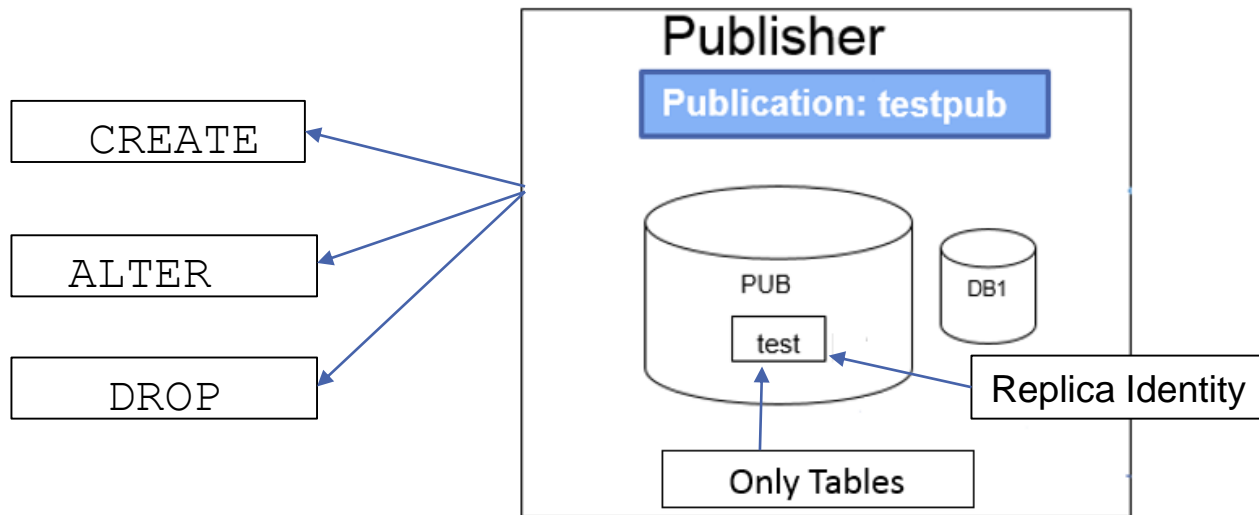


# Terminology used in Logical Replication:

---

- Publisher/Publication
- Subscriber/Subscription
- Replication Slot
- Replica Identity

Server1=pg10srv1



- First create tables to include in publication and then create publication.

- Publisher → node where a publication is defined.
- Publication → set of changes defined for tables.
- Publication can be defined on any ***physical replication master***.
- DB name can be different from Subscriber.
- Create publication using the 'CREATE' command and can be later ALTERED and DROPPED.
- A published table should exist before creating publication and must have a "replica identity" (can be primary key or index).
- Operations possible on tables– INSERT, UPDATE, and DELETE, combinations, or ALL(default)
- UPDATE or DELETE operations on publisher without replica identity will give error.
- Same replica identity has to be set at subscription side as well.
- Currently replicates only **tables**.



## ■ Create:

- `CREATE PUBLICATION testpub FOR TABLE tb1, tb2;`
- `CREATE PUBLICATION testpub FOR TABLE tbl3 WITH (publish = 'insert, update');`
- `CREATE PUBLICATION testpub FOR ALL TABLES;`

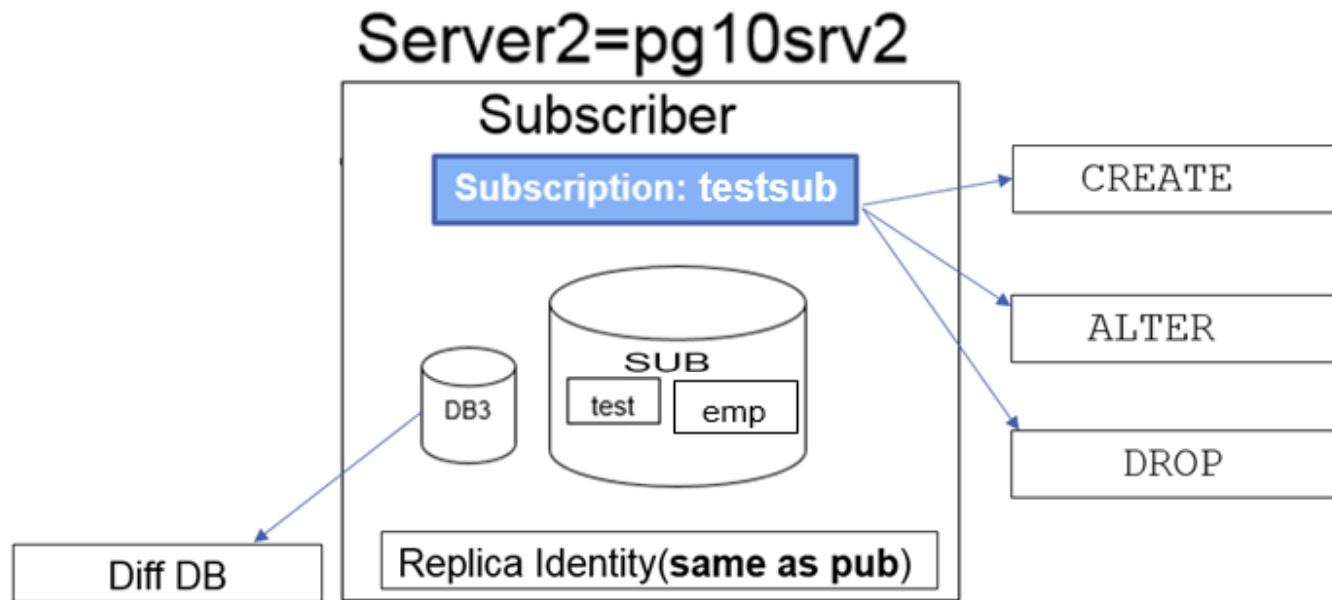
## ■ Alter :

- `ALTER PUBLICATION testpub ADD TABLE tb4, tbl5;`
- `ALTER PUBLICATION testpub SET TABLE tbl4, tbl5;`
- `ALTER PUBLICATION testpub DROP TABLE tbl4, tbl5;`

## ■ Drop:

- `DROP PUBLICATION testpub;`
- `DROP PUBLICATION testpub, testpub1;`

# Subscriber/Subscription



- First create tables with same name, same column name and type, same replica identity as publication.
- Later create subscription.

# Subscription:

- Subscription → Downstream side of logical replication and this node is called subscriber.
- It can be used as a publisher for other databases.
- Multiple subscriptions on one subscriber is possible.
- Subscription should be created using superuser.
- Same table names at publisher and subscriber.
- Tables should be created first before creating a subscription.
- Same columns in source and target tables with same data type but can have different order.

# Commands related to Subscription:

## ■ Create:

- `CREATE SUBSCRIPTION testsub1 CONNECTION 'host=<remote> 'dbname=<db> user=<user>' PUBLICATION testpub1,testpub2;`
- `CREATE SUBSCRIPTION mysub1 CONNECTION '...' PUBLICATION ...WITH (enabled = false, create_slot = false, slot_name = 'myslot', copy_data = false, ...);`

## ■ Alter:

- `ALTER SUBSCRIPTION testsub ENABLE/DISABLE;`
- `ALTER SUBSCRIPTION testsub CONNECTION 'host=newhost ...';`
- `ALTER SUBSCRIPTION testsub SET(slot_name = 'newslot'/'NONE');`
- `ALTER SUBSCRIPTION testsub REFRESH PUBLICATION;`

## ■ Drop:

- `DROP SUBSCRIPTION testsub;`

- For creating publication : `CREATE` privilege in the database.
- To add tables to a publication, the user must have `Ownership` rights on the table.
- To create a subscription, the user must be a superuser.
- Role used for the replication connection must have the `REPLICATION` attribute (or be a superuser)
- Entry for replication role in `pg_hba.conf`

## ■ Parameters: to be set in `postgresql.conf`

### @Publisher

- `wal_level = logical`
- `# max_replication_slots` → default value= 10
- `# max_wal_senders` → default value= 10

### @Subscribers

- `# max_replication_slots` → default value= 10
- `# max_logical_replication_workers` → default value= 4
- `# max_worker_processes`
- `# max_sync_workers_per_subscription` → default value= 8

\* Note: For remote connections, set- `listen_addresses` and `pg_hba.conf`

# Quick Setup and Demo

# Settings required for demo:

## 1. Setting postgresql.conf @Publisher

➤ `wal_level = logical`

- Note: Default values of other required parameters will suffice for this demo.

## 2. Use replication role `(CREATE ROLE replication WITH REPLICATION PASSWORD ' LOGIN;`

## 3. Setting pg\_hba.conf on both nodes.

➤ `host all postgres 0.0.0.0/0 SCRAM-SHA-256`

➤ `host all replication 0.0.0.0/0 SCRAM-SHA-256`


\* Note : The values here depend on your actual network configuration



# Overview of Demo:

Server1=pg10srv1

Server2=pg10srv2



# Creating database 'PUB' on Publisher:

```
postgres=# CREATE DATABASE PUB;
```

---

CREATE DATABASE

# Creating table 'TEST' on Publisher:

```
pub=# CREATE TABLE TEST (ID NUMERIC PRIMARY  
      KEY, NAME TEXT) ;
```

CREATE TABLE

# Insert some records in table 'TEST':

```
pub=# INSERT INTO TEST VALUES (1, 'TOM') ,  
      (2, 'DICK');
```

INSERT 0 2

# Select records of table 'TEST' :

```
pub=# select * from test;
```



```
id | name
---+-----
 1 | TOM
 2 | DICK
(2 rows)
```

# Let's see whether there is any Publication or not?

pub=# \dRp

List of publications

Name	Owner	All tables	Inserts	Updates	Deletes
-----+-----+-----+-----+-----+-----					
(0 rows)					

# Let's create a Publication on Publisher node:

```
pub=# CREATE PUBLICATION testpub FOR TABLE TEST;
```

# CREATE PUBLICATION

# Let's check the publication status:

pub=# \dRp

# List of publications

Name	Owner	All tables	Inserts	Updates	Deletes
-----+	-----+	-----+	-----+	-----+	-----
-					
testpub	postgres	f	t	t	t
(1 row)					

\* Default operation performed on tables is 'All' i.e. Insert, update and delete.

# This publication should now be added to the tables:

---

```
pub=# \d test
```



Table "public.test"

Column	Type	Collation	Nullable	Default
id	numeric		not null	
name	text			

Indexes:

"test\_pkey" PRIMARY KEY, btree (id)

Publications:

"testpub"

**Now let's set up subscription on subscriber**

# Let's create a new database 'SUB' on Subscriber:

```
postgres=# CREATE DATABASE SUB;
```

CREATE DATABASE

# Creating table 'TEST' on Subscriber:

```
sub=# CREATE TABLE TEST (NAME TEXT , ID NUMERIC  
      PRIMARY KEY) ;
```

- Order of columns can be different.

CREATE TABLE

# Let's see whether there is any Subscription or not?

sub=# \dRs

List of subscriptions

Name	Owner	Enabled	Publication
------	-------	---------	-------------

-----+	-----+	-----+	-----
--------	--------	--------	-------

(0 rows)



# Let's create a Subscription on Subscriber node:

```
sub=# CREATE SUBSCRIPTION testsub CONNECTION '  
user=postgres host=pg10srv1 dbname=pub ' PUBLICATION  
testpub;
```

NOTICE: created replication slot "testsub" on  
publisher

CREATE SUBSCRIPTION

# Let's check the subscription status:

---

sub=# \dRs

# List of subscriptions

Name	Owner	Enabled	Publication
-----+	-----+	-----+	-----
testsub	postgres	t	{testpub}
(1 row)			

# Let's see whether we got the data in table or not?

```
sub=# select * from test;
```

```
name | id
-----+-----
TOM  |  1
DICK |  2
(2 rows)
```

# Lets verify subscription log:

```
-bash-4.2$ tail -f postgresql-Tue.log
```

```
LOG:  worker process: logical replication worker for
subscription 49304 (PID 12053) exited with exit code 1
LOG:  logical replication apply worker for subscription
"testsub" has started
LOG:  logical replication table synchronization worker for
subscription "testsub", table "test" has started
LOG:  logical replication table synchronization worker for
subscription "testsub", table "test" has finished
```



# Now, let's try to insert some data on Publisher node:

```
pub=# INSERT INTO TEST VALUES (3, 'AFTER REPLICATION' );
```

INSERT 0 1

# Select data of 'Test' table on Subscriber:

```
sub=#  select * from test;
```

id		name
-----+-----		
1		TOM
2		DICK
3		AFTER REPLICATION

(3 rows)

# Write operations @Subscriber: Creating table 'EMP'

```
sub=# CREATE TABLE EMP (EMPID NUMERIC PRIMARY  
      KEY, EMPNAME TEXT) ;
```

CREATE TABLE

```
sub=# INSERT INTO EMP VALUES (1, 'LEE'), (2, 'TOM');
```

```
empid | empname
-----+-----
      1 | LEE
      2 | TOM
(2 rows)
```



## ■ On Publisher

- `pg_stat_replication`
- `pg_replication_slots`
- `pg_stat_activity`
- `\dRp`

## ■ On Subscriber

- `pg_stat_subscription`
- `pg_stat_activity`
- `\dRs`

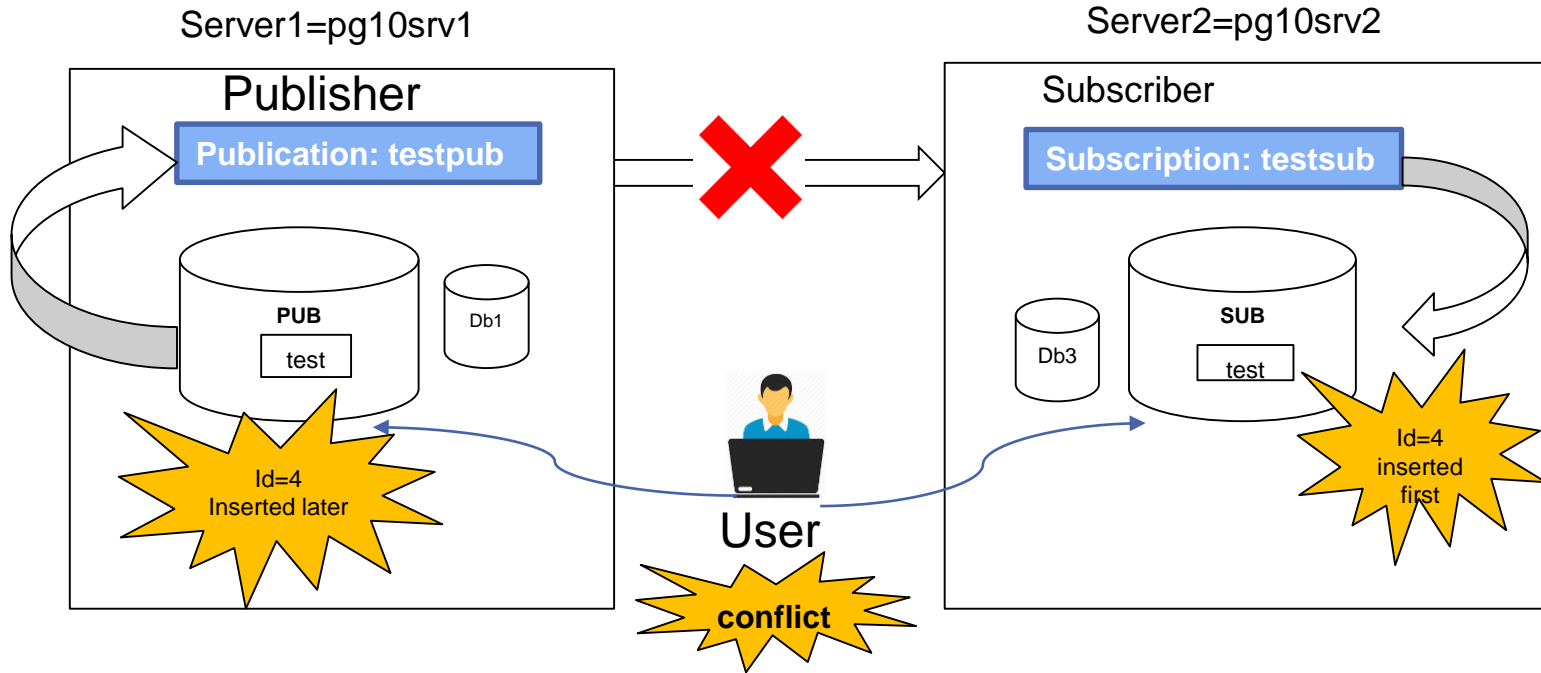
# Precautions at Subscriber:

- While Inserting rows.
- While modifying tables DDL – adding/dropping columns etc.
- If not, will leads to conflicts.

# CONFLICT

*If incoming data violates any constraints the replication will stop and is called CONFLICT*

# What if we insert a record with id=4 on Publisher?



# Check conflict status in subscriber log:

## On Subscriber

```
sub=# SELECT * FROM TEST;
```

name	id
TOM	1
DICK	2
AFTER REPLICATION	3
CONFLICT	4

(4 rows)

### Error in log file:

```
LOG: logical replication apply worker for subscription "testsub" has started  
ERROR: duplicate key value violates unique constraint "test_pkey"  
DETAIL: Key (id)=(4) already exists.
```

## On Publisher

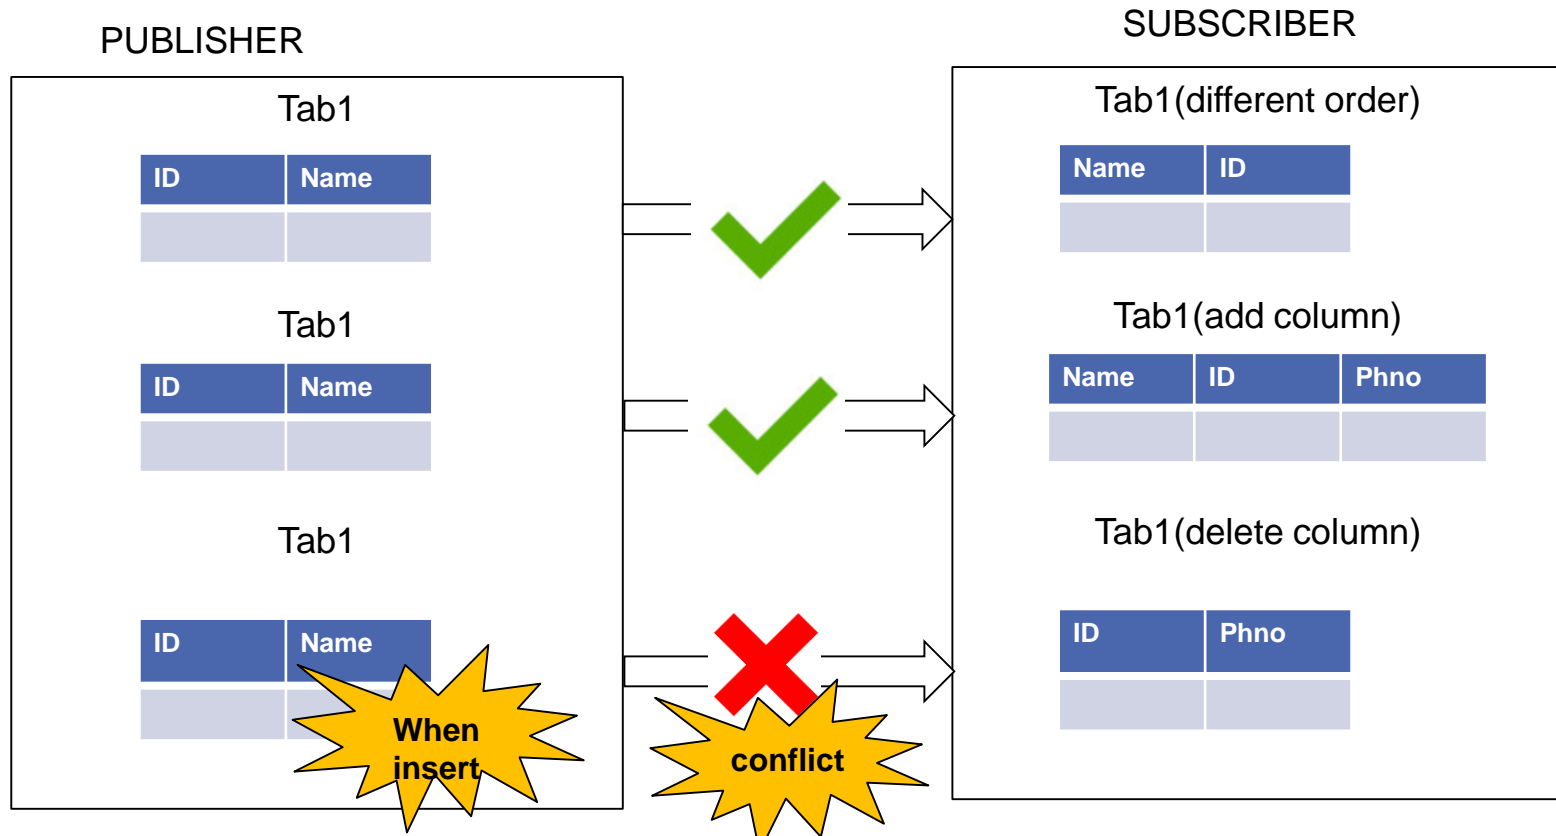
```
pub=# INSERT INTO TEST  
VALUES (4, 'RESOLVE');
```

```
sub=# SELECT * FROM  
TEST;
```

name	id
TOM	1
DICK	2
AFTER REPLICATION	3
CONFLICT	4

(4 rows)


# Delete a column on subscriber:



# Resolving Conflict:

- Change data on subscriber i.e. delete conflicting key.
- Skip the conflicting transaction from replication by calling the `pg_replication_origin_advance()` function.
- The current position of origins can be seen in the `pg_replication_origin_status` system view.

# Restrictions:



- Schema/DDL replication
- Sequences replication
- TRUNCATE command replication
- Large objects replication
- Views, Materialized views ,Partition root tables, or foreign tables replication



# Conclusion:

- Good feature but still room to include more features.
- No need to replicate whole instance.
- Subscriber can be used for Write operations but with care.
- Good for cross platform replication, different major version replication, write operations on subscriber, like use cases.
- Not the replacement of SR.

Thank you

# Your speaker



+61 2 9452 9017



[rajnib@fast.au.fujitsu.com](mailto:rajnib@fast.au.fujitsu.com)



[postgresql.fastware.com](http://postgresql.fastware.com)



[twitter.com/fujitsupostgres](https://twitter.com/fujitsupostgres)



[linkedin.com/showcase/fujitsu-enterprtise-postgres](https://linkedin.com/showcase/fujitsu-enterprtise-postgres)

## Rajni Baliyan

Database Administrator

Fujitsu Enterprise Postgres / PostgreSQL



# Q&A